

Pooja Swarnkar

Infrastructure & DevOps Engineer

Ahmedabad, Gujarat, India | +91 88246 75665 | poojasonii0206@gmail.com | linkedin.com/in/pooja-swarnkar | github.com/poojaswrnkar | poojaswrnkar.github.io

PROFESSIONAL SUMMARY

DevOps and Infrastructure Engineer with 2+ years of production experience across cloud architecture, CI/CD automation, and Kubernetes-based container orchestration, with hands-on ownership of on-premise AI infrastructure for large language model (LLM) inference. AWS Certified Cloud Practitioner. Track record includes standardizing deployment pipelines for ~20% faster releases, executing zero-downtime database and cache migrations, and architecting GPU-based model-serving infrastructure across multiple inference engines (vLLM, SGLang, llama.cpp). Experienced across AWS, GCP, DigitalOcean, and AceCloud, with strong grounding in Infrastructure as Code (Terraform), GitOps (FluxCD), observability (Prometheus, Grafana), and database administration (PostgreSQL, MongoDB, Redis).

TECHNICAL SKILLS

Cloud Platforms: AWS (ECS, EC2, S3, RDS, IAM, VPC, Lambda, SNS, CloudWatch, Secrets Manager), GCP, DigitalOcean, AceCloud

Containers & Orchestration: Docker, Kubernetes, Helm

CI/CD & GitOps: GitHub Actions, Jenkins, FluxCD, Git

Infrastructure as Code: Terraform

Observability & Reliability: Prometheus, Grafana, CloudWatch, High Availability, Incident Alerting

Databases: PostgreSQL, MongoDB, Redis, DragonflyDB

AI Infrastructure: LiteLLM, vLLM, SGLang, llama.cpp, LM Studio, GGUF, Quantization, GPU Infrastructure (NVIDIA RTX 6000 Ada, RTX PRO 6000 Blackwell, RTX A6000, L40S)

Security & Compliance: Secrets Management, SOC 2, Cloudflare (DDoS Protection)

Operating Systems: Linux, Bash, Networking, Nginx

PROFESSIONAL EXPERIENCE

DevOps Engineer

Sep 2025 - Present

Ankpal, Ahmedabad, India

- Own end-to-end environment lifecycle management across Dev, QA, UAT, and Production, enabling frequent, reliable software releases.
- Standardized CI/CD pipelines using GitHub Actions across all environments, reducing manual intervention and improving deployment speed by approximately 20%.
- Architected production infrastructure on AWS ECS, improving application scalability, availability, and uptime.
- Planned and executed a cross-account AWS RDS PostgreSQL migration using snapshot-based strategy, schema validation, consistency checks, and rehearsed rollback procedures, with zero downtime.
- Administer PostgreSQL (self-hosted and RDS), MongoDB, and Redis, covering high availability, automated backups, monitoring, and performance tuning.
- Deployed multiple Qwen model families (2.5 and 3.5, in 27B and 35B variants) for on-premise LLM inference across a multi-GPU fleet (NVIDIA RTX 6000 Ada, RTX PRO 6000 Blackwell, RTX A6000, L40S), selecting the inference engine per GPU VRAM capacity.
- Proxied the multi-model, multi-engine inference fleet through a single LiteLLM gateway with environment-scoped API keys and per-model uptime alerting.
- Collaborate closely with development and QA teams to ensure smooth, zero-defect production releases.

Freelance DevOps Engineer

Jun 2025 - Aug 2025

Independent - three concurrent client engagements

- Delivered end-to-end DevOps solutions for three concurrent client projects, focused on automation, security, and cost efficiency.
- HolidayTribe.ai: Managed Kubernetes deployments across production and non-production using Helm; implemented GitOps with FluxCD; built custom Jenkins pipelines for microservices; integrated AWS Secrets Manager with Terraform-managed infrastructure.

- NoticeBot (GST Litigation Platform): Deployed containerized frontend and backend on AWS EC2 using Docker; optimized image sizes for over 30% faster CI builds.
- Inphynitics: Deployed Java and NestJS backend services on AWS; hosted frontend on S3, improving availability and reducing hosting costs.

DevOps Engineer

Jun 2024 - Jun 2025

Linearloop.io, Ahmedabad, India

- Built and maintained automated CI/CD pipelines using GitHub Actions and Jenkins, reducing release cycle times across multiple products.
- Implemented Infrastructure as Code using Terraform on AWS for repeatable, version-controlled provisioning.
- Integrated Prometheus and Grafana for end-to-end observability, enabling proactive issue detection and alerting.
- Reduced cloud spend by auditing and optimizing resource utilization, scaling configurations, and eliminating waste.
- Spearheaded infrastructure for BetterBugs, a company product, from an empty AWS account: full ECS, Docker, and Terraform stack, Cloudflare DDoS protection, and SOC 2 compliance support.
- Migrated Redis to DragonflyDB with zero downtime; automated AWS alerts via Lambda and SNS.

DevOps Intern

Jan 2024 - Mar 2024

Sudowall, Udaipur, India

- Configured GitHub-Jenkins integration for automated build and deployment pipelines.
- Deployed a Prometheus and Grafana monitoring stack using Docker Compose for real-time infrastructure visibility.
- Gained hands-on experience with core AWS services: IAM, VPC, EC2, S3, SNS, and CloudWatch.

KEY PROJECTS

BetterBugs - Bug Tracking SaaS (app.betterbugs.io)

Founding Infrastructure Engineer, via Linearloop.io

- Designed and built the entire cloud infrastructure from scratch using AWS ECS, Docker, and Terraform with full IaC coverage.
- Achieved a zero-downtime Redis to DragonflyDB migration; secured infrastructure with Cloudflare DDoS protection and AWS Secrets Manager.
- Implemented Prometheus and Grafana monitoring with automated Lambda and SNS alerting; supported SOC 2 compliance evidence collection.

On-Premise LLM Infrastructure

Ankpal - internal AI infrastructure product

- Deployed and benchmarked multiple Qwen model families (2.5, 3.5, 3.6) across a four-GPU fleet: RTX 6000 Ada, RTX PRO 6000 Blackwell, RTX A6000, and L40S.
- Selected inference engine per GPU VRAM capacity: llama.cpp, via LM Studio, on 48GB-VRAM cards, and vLLM on 96GB-VRAM cards, with SGLang benchmarked alongside both.
- Unified all models behind a single LiteLLM gateway with environment-scoped API keys and per-model health alerting.

EDUCATION

Bachelor of Technology (B.Tech), Computer Science

2024

Geetanjali Institute of Technology, Udaipur | CGPA: 9.1 / 10

CERTIFICATIONS

AWS Certified Cloud Practitioner - Amazon Web Services

cp.certmetrics.com/amazon/en/public/verify/credential/4440f705f7e44148acd17289cd49eac1